



Exercise Sheet 9

Intelligent Systems

January 06, 2020

Anomalies and Classification

Exercise 1 - DBSCAN and Outlier Detection

- A. Calculate the *Local Outlier Factor (LOF)* of the points A_1 and N in Figure 1.
- B. Draw the distribution of ascending minimum *kdist*s

$$\min_y(kdist(x, y)),$$

of every point with $k = 1, 2, 3$.

- C. How can you estimate the parameter ϵ by given a percentage of noise?
- D. Find parameters $\epsilon > 0$, $\min_pts \in \mathbb{N}$ s.t.
 - $A_i, i = 1, 2, 3$ is clustered as a cluster
 - $B_j, j = 1, 2$ is clustered as a cluster
 - N is marked as noise.
- E. Find parameters $\epsilon > 0$, $\min_pts \in \mathbb{N}$, and points C_k s.t.
 - $A_i, i = 1, 2, 3, B_j, j = 1, 2$ is clustered as a cluster.
 - N is marked as noise.

Exercise 2 - Classification algorithms

- A. Observe the data set in Table 1. First, create a 1-R Classifier that is able to predict whether a person is going to visit the party this evening by using the information of his/her amount of money, whether he/she writes an exam tomorrow, or if his/her heartthrob will come to the party.
- B. Extend your 1-R Classifier to a Decision Tree. Which features should be placed on higher levels of the tree?
- C. Apply the Naïve Bayes Classifier on the same data set. Calculate also the probabilities $P(Yes|E1)$ and $P(No|E6)$.




Figure 1: Points.

| Sample | Money | Exam | Heartthrob | Party |
|--------|-------|------|------------|-------|
| E1 | 10 | Yes | Yes | Yes |
| E2 | 13 | No | Yes | Yes |
| E3 | 11 | Yes | No | No |
| E4 | 12 | No | No | Yes |
| E5 | 7 | Yes | Yes | Yes |
| E6 | 5 | Yes | No | No |
| E7 | 6 | No | Yes | Yes |
| E8 | 8 | No | No | No |

Table 1: Party Datensatz

Exercise 3 - Classification with Python

Download the Jupyter Notebook *7_Classification.ipynb* and the dataset *usStatesData2.csv* from *OpenOlat*. The task is to predict the vote label only by given the features in the dataset (e.g. *percentage of weapon owners*). Which classification metrics would you apply to test the performance of your classifier?

Signature Task - Classification Challenge

In classification tasks, the labels (here: the author of the signature) in the training data is given to the model for training. The task is to build a classifier that is able to predict the author of a signature. To present your results make a short presentation with 4 slides describing the following steps:

- Preprocessing (Outlier detection/Normalisation/Segmentation)
- Extraction of your lower dimensional features
- Dimensionality reduction via PCA
- Classification (Bayes Classifier, Decision Trees, SVM, NN, etc.)